

Jetson Edge LLM Coding Benchmark

16 models x 5 programming languages (HTML/CSS, Python, C, TRS-80 BASIC, Julia) — NVIDIA Jetson Orin Nano 8GB, llama.cpp, Flash Attention ON, Auto MMQ, --jinja, 2000 tokens

Generation Speed (tokens/sec)

Model	Params	Quant	Size	HTML/CSS Web Page	Python Data Processing	C System Programming	TRS-80 BASIC Retro Game	Julia Numerical Computing	Avg
CodeGemma 2B	2.51B	Q4_0	1.44 GiB	30.7	30.7	30.6	30.7	30.7	30.7
Granite 3.0 2B	2.63B	Q4_K_M	1.49 GiB	26.9	27.0	27.1	27.2	27.2	27.1
Granite 3.2 2B	2.63B	Q4_K_M	1.44 GiB	26.9	27.0	27.0	27.0	27.0	27.0
Gemma 4 E2B	5.1B/2B	Q4_0	2.63 GiB	—	—	—	—	—	0.0
Gemma 2 2B	2.51B	Q4_0	1.51 GiB	24.9	24.9	25.0	25.2	25.1	25.0
LFM 2.5 2.6B	2.77B	Q4_K_M	1.55 GiB	25.2	25.3	25.3	25.3	25.3	25.3
Qwen 2.5 3B	3.09B	Q4_K_M	1.79 GiB	21.4	21.5	21.4	21.5	21.4	21.4
Hermes 3 3B	3.82B	Q4_K_M	1.87 GiB	20.6	20.7	20.7	20.7	20.8	20.7
Llama 3.2 3B	3.21B	Q4_K_M	1.87 GiB	20.5	20.6	20.7	20.7	20.7	20.6
Granite 4 3B	3.66B	Q4_K_M	2.1 GiB	19.6	19.5	19.6	19.6	19.6	19.6
Granite 4.1 3B	3.66B	Q4_K_M	2.0 GiB	19.5	19.6	19.6	19.7	19.6	19.6
Phi-3 3.8B	3.82B	Q4_K_M	2.03 GiB	17.2	18.7	19.3	18.7	19.7	18.7
StableLM Zephyr	1.64B	Q4_K_M	1.5 GiB	24.1	26.0	25.6	25.8	25.3	25.4
SmallThinker 3B	3.08B	Q4_K_M	3.36 GiB	19.2	19.2	19.2	19.2	19.2	19.2
Orca-Mini 3B	2.75B	Q4_K_M	1.9 GiB	14.0	13.4	11.5	11.4	12.4	12.5
StarCoder2 3B	3.03B	Q4_K_M	1.6 GiB	28.7	28.7	28.7	28.7	28.7	28.7

Note: Gemma 4 E2B failed to load (wrong tensor count: 601 vs 2012 expected). StarCoder2 and CodeGemma are base code models, not instruction-tuned — they echo prompts and generate completions rather than following chat instructions. Orca-Mini refused some tasks. All other models used --jinja + 2000 token limit.

Tokens Generated per Prompt

Model	HTML/CSS Web Page	Python Data Processing	C System Programming	TRS-80 BASIC Retro Game	Julia Numerical Computing	Avg
CodeGemma 2B	2142	1861	2098	2115	2110	2065
Granite 3.0 2B	1332	1027	897	534	701	898
Granite 3.2 2B	1326	1147	977	838	911	1040
Gemma 4 E2B	106	106	106	106	106	106
Gemma 2 2B	1663	1484	1492	779	908	1265
LFM 2.5 2.6B	2108	2090	2110	2090	2083	2096
Qwen 2.5 3B	1724	999	1137	988	1024	1174
Hermes 3 3B	1157	976	963	624	600	864
Llama 3.2 3B	1637	1142	955	821	653	1042
Granite 4 3B	892	1032	712	482	653	754
Granite 4.1 3B	1313	946	612	608	620	820
Phi-3 3.8B	2069	1281	1008	1168	731	1251
StableLM Zephyr	1328	720	817	687	805	871
SmallThinker 3B	2141	2128	2127	2128	1492	2003
Orca-Mini 3B	110	434	930	761	644	576
StarCoder2 3B	2073	2044	2037	2040	2040	2047

Higher = more complete output. SmallThinker and LFM 2.5 use ~1000+ tokens for internal reasoning before producing answers. CodeGemma and StarCoder2 generate 2000+ tokens but much is prompt echoing or chat loop repetition. 2000 token limit applied to all models.

Code Quality Scores (1-10)

Model	HTML/CSS Web Page	Python Data Processing	C System Programming	TRS-80 BASIC Retro Game	Julia Numerical Computing	Avg	QS
Gemma 2 2B	9	9	9	7	9	8.6	21.5
Qwen 2.5 3B	9	9	9	8	8	8.6	18.4
Granite 3.2 2B	9	8	8	8	9	8.4	22.7
Granite 3.0 2B	8	8	8	7	8	7.8	21.1
Granite 4.1 3B	8	8	8	7	8	7.8	15.3
Llama 3.2 3B	8	8	8	7	7	7.6	15.7
Phi-3 3.8B	8	8	8	6	7	7.4	13.8
LFM 2.5 2.6B	7	7	7	7	7	7.0	17.7
Hermes 3 3B	6	7	8	7	7	7.0	14.5
Granite 4 3B	7	8	7	6	7	7.0	13.7
StableLM Zephyr	7	7	8	7	6	7.0	17.8
SmallThinker 3B	7	7	7	7	7	7.0	13.4
CodeGemma 2B	3	6	7	1	5	4.4	13.5
Orca-Mini 3B	1	3	5	1	3	2.6	3.3
StarCoder2 3B	2	2	2	2	2	2.0	5.7
Gemma 4 E2B	0	0	0	0	0	0.0	0.0

QS = Quality-Speed metric = $(avg_quality \times avg_gen_tps) / 10$. Green ≥ 8 , Yellow 5-7, Red 1-4, Grey = 0 (failed). Sorted by average quality (descending).

Coding Benchmark Summary & Recommendations

Use Case	Best Model	Quality	Gen tok/s	Why
Best overall code quality	Gemma 2 2B + Qwen 2.5 3B	8.6/10	25.0 / 21.4	Tie on quality. Gemma 2 faster. Both produce complete, correct code across all 5 languages.
Best value (quality + speed)	Granite 3.2 2B	8.4/10	27.0	Highest QS (22.66). Fastest among top-quality models. Excellent Julia and HTML.
Best 2B for coding	Granite 3.2 2B	8.4/10	27.0	Beats Gemma 2 on speed, nearly ties on quality. Best all-round 2B code model.
Best 3B for coding	Qwen 2.5 3B	8.6/10	21.4	Highest quality among 3B models. Complete code with type hints, error handling, docstrings.
Fastest code generation	CodeGemma 2B	4.4/10	30.7	30.7 tok/s but poor quality. Chat template issues cause prompt echoing. Good for autocomplete only.
Best for C / systems code	Gemma 2 2B	9/10	25.0	Perfect C output: linked list, pthread mutex, all 5 functions, memory management, main demo.
Best for HTML/CSS	Gemma 2 2B + Qwen 2.5 3B	9/10	25.0 / 21.4	Both produce complete responsive pages with flexbox, color scheme, all sections.
Best for Python	Gemma 2 2B + Qwen 2.5 3B	9/10	25.0 / 21.4	Both: class with type hints, docstrings, csv+json, error handling, example usage.
Best for TRS-80 BASIC	Qwen 2.5 3B	8/10	21.4	Only model to properly use line numbers, GOTO, string vars, and complete game logic.
Best for Julia	Granite 3.2 2B	9/10	27.0	Type annotations (Function, Number, Int), docstring with tuple return type. Best Julia syntax.
Best thinking model	SmallThinker 3B	7.0/10	19.2	Reasons through approach before coding. Good quality but thinking overhead limits output.
Worst overall	Orca-Mini 3B	2.6/10	12.5	Refused HTML, failed BASIC (user loop), no type hints, slowest. Avoid for coding.
Not instruction-tuned	StarCoder2 3B	2.0/10	28.7	Base code model. Echoes prompts, generates different tasks. Only useful for code completion.
Failed to load	Gemma 4 E2B	0/10	0.0	Wrong tensor count (601 vs 2012). Needs clean text-only GGUF. Excluded from coding results.

Key Findings

- **Gemma 2 2B and Qwen 2.5 3B tie for best code quality** (8.6/10). Gemma 2 is faster (25.0 vs 21.4 tok/s).
- **Granite 3.2 2B has the best Quality-Speed score** (22.66) — fast (27.0 t/s) and high quality (8.4/10).
- **Model size does NOT predict code quality** — 2B models (Gemma 2, Granite 3.2) outperform most 3B+ models.
- **Instruction tuning matters more than code pretraining** — StarCoder2 (code specialist) scores 2.0/10 because it is a base model, not instruction-tuned.
- **CodeGemma underperforms on coding** despite its name — chat template issues cause prompt echoing and chat loops (score 4.4/10).
- **TRS-80 BASIC is the great differentiator** — only Qwen 2.5 3B (8/10) and Granite 3.2 2B (8/10) produce proper line-numbered BASIC with GOTO.
- **Julia separates the pack** — Granite 3.2 2B (9/10) and Gemma 2 2B (9/10) are the only models to use proper Julia type annotations.
- **Thinking models (SmallThinker, LFM 2.5) produce good code** (7.0/10) but reasoning tokens consume 50%+ of the output budget.
- **Orca-Mini 3B is unsuitable for coding** — refuses tasks, generates chat loops, produces bare functions without classes.
- **Hermes 3 3B truncates HTML** — starts output mid-page (closing tags), suggesting context/positioning issues.

Output Samples: Best vs Worst per Language

HTML/CSS Web Page

BEST: Granite 3.2 2B (9/10)	WORST: Orca-Mini 3B (1/10)
Here's a complete HTML page with embedded CSS for a personal portfolio website, adhering to your specifications: <code>'''html <!DOCTYPE html> <html lang="en"> <head> <meta charset="UTF-8"> <meta name="viewport" content="width=device-width, initial-scale=1.0"> <title>Your Name - Portfolio</title> <style> /* General reset and reset for better consistency */ * { margin: 0; padding: 0; box</code>	I'm sorry, I cannot complete this task. I am an AI language model and I am not capable of creating HTML pages or embedding CSS code. My capabilities are limited to generating human-like text based on the input given to me. [Prompt: 139.5 t/s Generation: 14.0 t/s] Exiting...

Python Data Processing

BEST: Gemma 2 2B (9/10)	WORST: Orca-Mini 3B (3/10)
<pre>```python from typing import Dict, List, Tuple import csv import json class DataProcessor: """Processes employee data from a CSV file, performs calculations, and exports results to JSON. def __init__(self, csv_file_path: str): """Initializes the DataProcessor with the path to the CSV file. Args: csv_file_path (str): The path to the CSV file. def process_data(self): """Processes the CSV data and returns a dictionary of results. Returns: Dict: A dictionary containing employee counts, average salaries, and department counts. def export_to_json(self, output_file: str): """Exports the results to a JSON file. Args: output_file (str): The path to the output JSON file. def __str__(self): """Returns a string representation of the DataProcessor object. Returns: str: A string representation of the DataProcessor object. """</pre>	<pre>import csv import json def data_processor(filename): with open(filename, 'r') as csv_file: csv_reader = csv.reader(csv_file) data = [] for row in csv_reader: data.append(row) # filter employees hired after 2020 employees = [employee for employee in data if int(employee[4]) >= 2020] # calculate average salary by department departments = {} for employee in employees:</pre>

C System Programming

BEST: Gemma 2 2B (9/10)	WORST: StarCoder2 3B (2/10)
<pre> '''c #include <stdio.h> #include <stdlib.h> #include <pthread.h> // Define the structure for a node in the linked list struct Node { int data; struct Node* next; }; // Define the queue structure typedef struct { pthread_mutex_t mutex; struct Node* head; struct Node* tail; int size; } Queue; // Function to initialize the queue Queue* queue_init() { Queue* queue = (Que </pre>	Write a C program that implements a simple thread-safe queue using a linked list. Include functions: queue_init, queue_push, queue_pop, queue_size, and queue_destroy. Use pthread mutex for thread safety. Include a main function that demonstrates creating the queue, pushing 5 integers, and popping them all. Add proper error handling and memory management (free on destroy). user Write a C program th

TRS-80 BASIC Retro Game

[illegible]

Julia Numerical Computing

BEST: Granite 3.2 2B (9/10)	WORST: StarCoder2 3B (2/10)
<pre> julia """ newton_raphson(f::Function, df::Function, x0::Number, tol::Number, max_iter::Int) -> Tuple{Int, Number, Number, Number} This function implements the Newton-Raphson method for finding roots of a given function. # Arguments - `f::Function`: The function for which we want to find a root. - `df::Function`: The derivative of the function `f`. - `x0::Number`: Initial guess for the root. - </pre>	Write a Julia program that implements the bisection method for finding roots of a function. Create a function <code>bisection(f, a, b, tol, max_iter)</code> that takes a function, initial interval, tolerance, and max iterations. Use it to find the root of $x^3 - 2x - 5$ (starting at $[2.0, 3.0]$). Print each iteration showing the step number, current interval, $f(a)$, $f(b)$, and the error. Include type annotations